PROGRAMMING FPGAS WITH DPC++

Adrian Jackson

a.jackson@epcc.ed.ac.uk

THE UNIVERSITY of EDINBURGH

epcc

What is an FPGA?

- Field programmable gate array
  - Programmable logic blocks
  - Memory units
  - Digital signal processing blocks
  - Interconnect
  - Other specialised pre-configured units (i.e. soft cores)
- Spatial architecture rather than instruction set architecture
  - Basic building blocks configurable to create your own processor
  - Reprogrammable as required
- Specialised processors design for individual applications rather than general purpose processor

THE UNIVERSITY of EDINBURGH

epcc

Instruction set architecture

- Processor implements pre-defined instruction set
  - Implementation provides functional units for each (set of) instruction(s)
  - Implementation defines fix control flow for each (set of instruction(s)
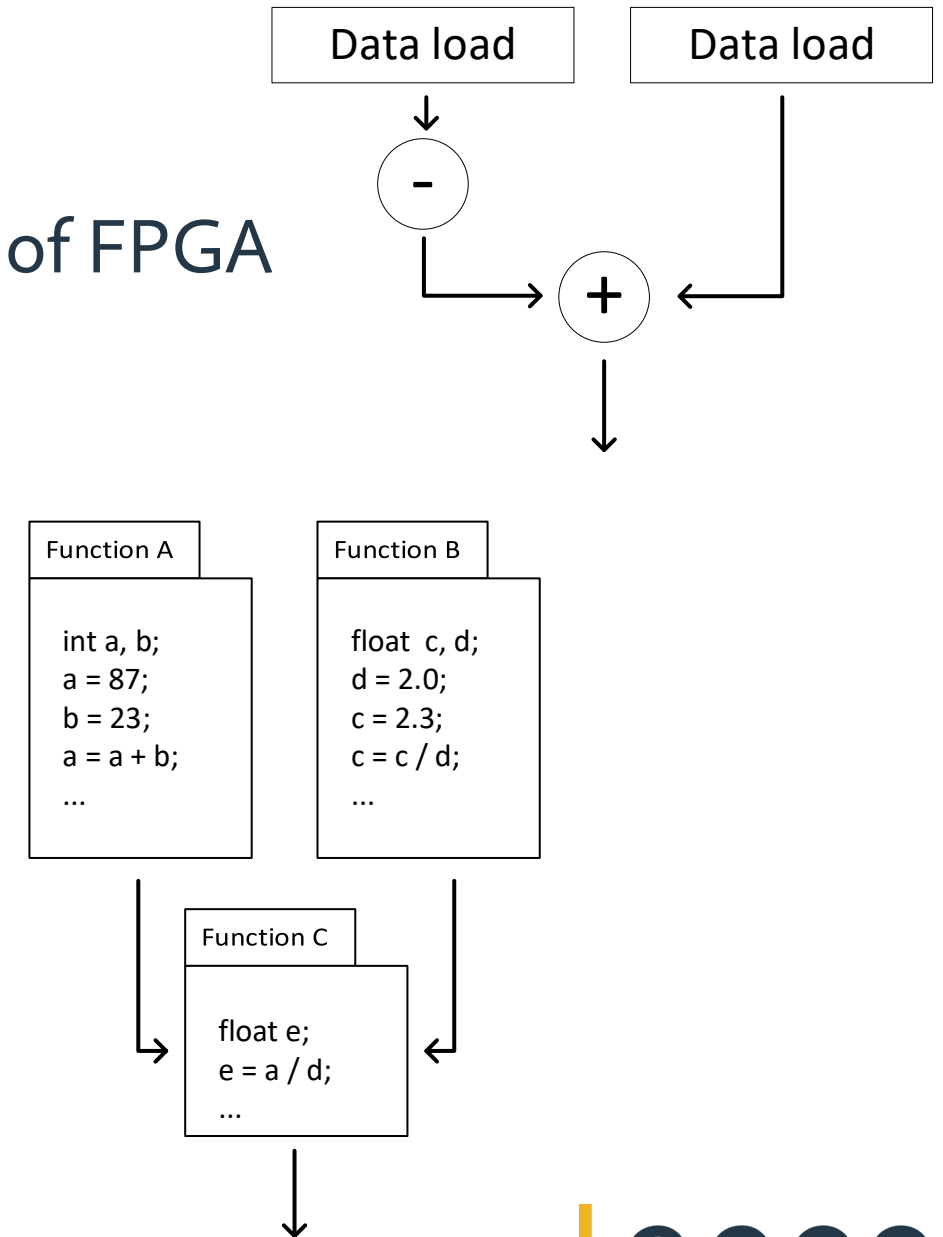 - Instruction flow architecture
  - Pipelining of instructions

Image from http://alanclements.org/simple%20isa.html

Dataflow architecture

- Individual program region encoded as part of FPGA

- Program region connected together

- Data flows through the hardware
  - Can recirculate as required

- All parts of the FPGA can be active

- Programmed on run
  - Hardware configuration pre-built

- Spatial computing


```
Function A

int a, b;
a = 87;
b = 23;
a = a + b;
...
```

```
Function B

float  c, d;
d = 2.0;
c = 2.3;
c = c / d;
...
```

```
Function C

float e;
e = a / d;
...
```

Why FPGAs?

- Processor design tools

- Enable removing unused hardware components
  - Enable 100% usage of hardware

- Sidestep ISA limitations
  - i.e. dependency chains for vectorization
  - i.e. instruction issue bottlenecks
  - i.e. register pressure
  - i.e. fixed width precisions
  - etc...

- High memory bandwidth and data connectivity
  - Streaming data processing etc...

THE UNIVERSITY of EDINBURGH

| epcc |

Why not FPGAs?

- Programming is hard
  - Low level
  - Multiple dimensional optimisation problem

- Portability isn't straight forward

- Large program performance can be limited
  - FPGA footprint is key to efficiency
  - Reuse of space can fit large programs onto the device

- Fixed hardware blocks need to be used to get maximum performance

- Dynamic reprogramming possible, but dangerous (glitch)

- Programming can only be done when powered on
 - But program maintained once loaded

THE UNIVERSITY of EDINBURGH

epcc

FPGA components

- ## Programmable logical blocks ■
  - ### Lookup tables
  - ### Adders to combine
  - ### Mimic hard gates

Truth Table

Inputs				Output
A	B	C	D	Y
0	0	0	0	0
0	0	0	1	0
0	0	1	0	0
0	0	1	1	1
0	1	0	0	0
0	1	0	1	1
0	1	1	0	1
0	1	1	1	0
1	0	0	0	0
1	0	0	1	1
1	0	1	0	1
1	0	1	1	0
1	1	0	0	1
1	1	0	1	0
1	1	1	0	0
1	1	1	1	0

epcc

FPGA components

- ## Local memory units
  - Block RAM
  - Multiple connections (Ports) possible
  - Configurable widths

- ## DSP blocks
  - Constrained FPUs
  - Multiplexers and ALUs

THE UNIVERSITY of EDINBURGH

epcc

FPGA components

- Interconnects
  - Routing connections between different components
  - Configurable
  - High overall bandwidth

- Connection blocks
  - PCI connections
  - Memory controllers

- Hard cores
  - Embedded full hardware implementations
  - i.e. Intel Stratix has:
    - Quad-core 64-bit ARM Cortex-A53 processor
    - PCIe Gen1/Gen2/Gen3 complete protocol stack
    - 10/25/100 Gbps Ethernet MAC
    - DDR4/DDR3/LPDDR3 hard memory controller

Programming FPGAs

- Traditional programming
  - Map program to ISA: Compiler
  - Language/compiler specification maps language features to ISAs
  - Optimisation done on specific ISA constraints and specific processor constraints

- FPGA programming
  - Deciding how to configure the hardware for the program you want to run
  - Decide what component implementations to put where on the FPGA (placement and routing)
  - Connect up all the external resources (memory and networks)
  - Create FPGA bitstream to program the device

Traditional FPGA programming

- Low level design tools
  - VHDL, Verilog
  - Write LUTs for specific hardware
  - Create layout and connection maps
  - Synthesis
  - Placement
  - Routing


```vhdl
1 library ieee;
2 use ieee.std_logic_1164.all;
3 use ieee.numeric_std.all;
4
5 entity signed_adder is
6   port
7   (
8     aclr : in      std_logic;
9     clk  : in      std_logic;
10    a    : in      std_logic_vector;
11    b    : in      std_logic_vector;
12    q    : out     std_logic_vector
13  );
14 end signed_adder;
15
16 architecture signed_adder_arch of signed_adder is
17   signal q_s : signed(a'high+1 downto 0); -- extra bit wide
18
19 begin -- architecture
20   assert (a'length >= b'length)
21     report "Port A must be the longer vector if different sizes!"
22     severity FAILURE;
23   q <= std_logic_vector(q_s);
24
25   adding_proc:
26   process (aclr, clk)
27     begin
28       if (aclr = '1') then
29         q_s <= (others => '0');
30       elsif rising_edge(clk) then
31         q_s <= ('0'&signed(a)) + ('0'&signed(b));
32       end if; -- clk'd
33     end process;
34
35 end signed_adder_arch;
```

```verilog
module tb ();
  reg a, b, c;
  wire y;

  design dut (.a(a), .b(b), .c(c), y(y)

  // apply input sequence

  initial
    begin

      a = 0; b = 0; c = 0 ; # 10;
      a = 0; b = 0; c = 1 ; # 10;
      a = 0; b = 1; c = 0 ; # 10;
      a = 0; b = 1; c = 1 ; # 10;
      a = 1; b = 0; c = 0 ; # 10;
      a = 1; b = 0; c = 1 ; # 10;
      a = 1; b = 1; c = 0 ; # 10;
      a = 1; b = 1; c = 1 ; # 10;

    end
endmodule
```

Higher level programming

- ## C to Gates
  - C to Register Transfer Level common tool for productivity

- ## Automated tools to take higher level languages to bitstream
  - Use compiler to convert program into LUTs and DSP blocks
  - Undertake placement, routing, and synthesis
  - Highly unconstrained problem
  - Still requires user help and input

- ## OpenCL took this approach
  - Basis for other high-level tools

epcc
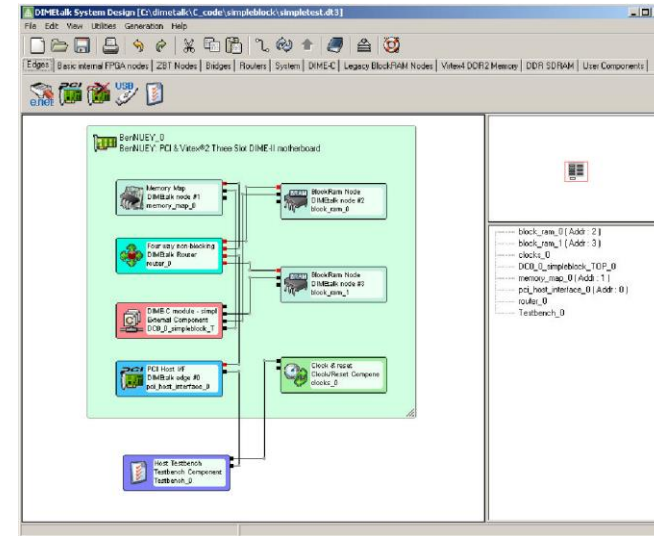
OneAPI DPC++

- DPC++ follows on from OpenCL
  - Uses SYCL to port to FPGA
  - Integrates Intel's (Altera's) Quartus FPGA tools
- Convert application code into SYCL
- Convert SYCL to HDL using HLS
- Synthesis HLS and build bitstream
- Run the code and program the FPGA at runtime

```
#include <CL/sycl.hpp>
using namespace sycl;

static const int N = 8;

int main(){
  queue q;
  std::cout << "Device: " <<
q.get_device().get_info<info::device::name>() <<
std::endl;

  int *data = malloc_shared<int>(N, q);
  for(int i=0; i<N; i++) data[i] = i;

  q.parallel_for(range<1>(N), [=] (id<1> i){
    data[i] *= 2;
  }).wait();

  for(int i=0; i<N; i++) std::cout << data[i] <<
std::endl;
  free(data, q);
  return 0;
}
```

THE UNIVERSITY *of* EDINBURGH

epcc

Compile cycle

- ## Synthesis (routing and placement) is slow
  - ### Hours

- ## Simulated/emulate synthesis can be used to test build and compile phase
  - ### Enables quick compile/development loop
  - ### Does not necessarily pick up all hardware issues

Device Image Type	Time to Compile	Description
FPGA Emulator	Seconds	The FPGA device code is compiled to the CPU. Use the Intel® FPGA Emulation Platform for OpenCL™ software to verify your SYCL code's functional correctness.
FPGA Simulator	Minutes	The FPGA device code is compiled to the CPU. Use the Questa*-Intel® FPGA Edition simulator to debug your code.
Optimization Report	Minutes	The FPGA device code is partially compiled for hardware. The compiler generates an optimization report that describes the structures generated on the FPGA, identifies performance bottlenecks, and estimates resource utilization.
FPGA Hardware Image	Hours	Generates the real FPGA bitstream to execute on the target FPGA platform.

epcc

Compilation

- ## DPC++ standard tools

  ```
  dpcpp –fintelfpga -Xshardware my_source_code.cpp
  dpcpp –fintelfpga -Xsboard=intel_a10gx_pac my_source_code.cpp
  ```

- ## Emulate rather than compile fully

  ```
  dpcpp –fintelfpga my_source_code.cpp
  ```

- ## Optimisation report

  ```
  dpcpp -fsycl-link=early my_source_code.cpp
  ```

- ## Simulate rather than emulate

  ```
  dpcpp -fintelfpga fpga_compile.cpp -Xssimulation -Xsboard=intel_s10sx_pac:pac_s10
  ```

THE UNIVERSITY of EDINBURGH

epcc

Program design

- Pipeline parallelism is the key approach
  - Aim for all hardware parts to be active all the time (after startup)
  - High throughput (reduced clock speed)
  - Large aggregate memory bandwidth
 - External memory access can reduce performance

Pipelined Spatial Compute

Pipeline stage 1	Memory Load
Pipeline stage 2	* — 10
Pipeline stage 3	+ — 42
Pipeline stage 4	Memory Store

Reinders, J., Ashbaugh, B., Brodman, J., Kinsner, M., Pennycook, J., Tian, X. (2021). Programming for FPGAs. In: Data Parallel C++.
https://link.springer.com/content/pdf/10.1007/978-1-4842-5574-2.pdf

THE UNIVERSITY of EDINBURGH

epcc

Program design

- ## Spatial parallelism
  - Program blocks mapped to hardware

- ## Memory/data parallelism
  - Direct memory configuration

- ## Both loop and vectorisation type optimisation possible
  - FPGA is often considered for spatial parallelism, but loop or vectorisation styles can also be mapped to the FPGA
  - Vectorisation may not provide the most efficient method (high space usage)
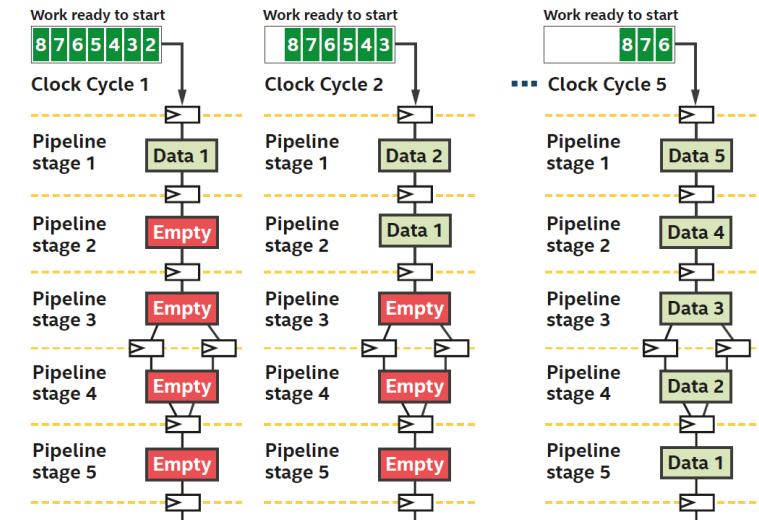

```
int a = 0;
for (int i=0; i < size; i++) {
    a = a + i;
}
```

epcc
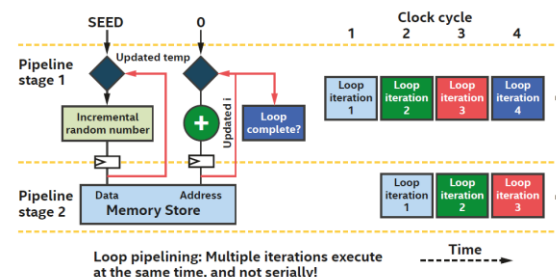
Program design

- ## Instruction level parallelism
  - ### Done at compile time
  - ### Trade off with space usage

- ## Vectorisation
  - ### Old school pipeline vectorization

- ## Loop unrolling
  - ### Spread loop iterations across hardware

- ## Conditional unbundling

```
c = a + b;
f = d * e;
```

```
LV    v1
MULV  v3,v1,v2
ADDV  v5, v3, v4
```

```
int result;
if (x > 0)
    result = foo (a);
else
    result = foo (b);
```

```
int result;
/* foo(a) and foo(b) both execute */
int ra = foo(a);
int rb = foo(b);
/* result is multiplexed */
if (x > 0)
    result = ra;
else
    result = rb;
```
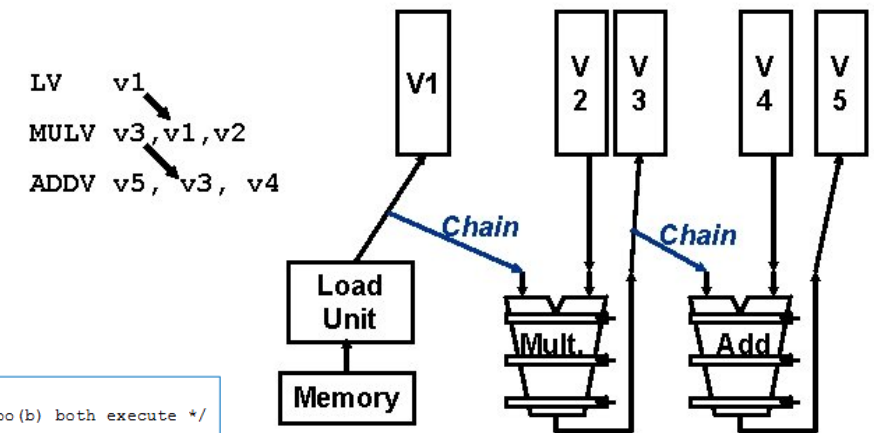
Program design

- Exposing parallelism is wider than other systems
  - Functions
  - Loops
  - Vectors
  - Whole call stacks
  - etc...

- Data costs still as important
  - Reuse
  - Locality
  - etc...

DPC++

- Keeping pipelines full key for performance

- ND-range kernels allow compiler to queue and dispatch these

```
h.parallel_for({16,16,16}, [=] (auto I) {
    output[I] = generate_random_number_from_ID(I);
});
```

- Standard loops allow iterations to be queued and dispatched
  - Unrolling allows further parallelisation
  - Space trade off vs throughput/pipeline depth

```
#pragma unroll 4
for(i = 0 ; i < 20; i++){
    a[i] += 1;
}
```

THE UNIVERSITY of EDINBURGH

epcc

Dependencies

```cpp
#include <CL/sycl.hpp>
using namespace sycl;
static const int N = 4;
int main(){
  queue q;
  int *d1 = malloc_shared<int>(N, q);
  int *d2 = malloc_shared<int>(N, q);
  for(int i=0; i<N; i++) { d1[i] = 10; d2[i] = 10; }

  auto e1 = q.parallel_for(range<1>(N), [=] (id<1> i){
    d1[i] += 2;
  });

  auto e2 = q.parallel_for(range<1>(N), [=] (id<1> i){
    d2[i] += 3;
  });

  q.parallel_for(range<1>(N), {e1, e2},  [=] (id<1> i){
    d1[i] += d2[i];
  }).wait();

  for(int i=0; i<N; i++) std::cout << d1[i] << std::endl;

  free(d1, q);
  free(d2, q);
  return 0;
}
```

DPC++

- Memory types
  - Functions for different memory functionality
    ```
    group_local_memory_for_overwrite()
    struct State {
      [[intel::fpga_memory]] int array[100];
      [[intel::fpga_register]] int reg[4];
    };
    cgh.single_task<class test>([=] {
      struct State S1;
      [[intel::fpga_memory]] struct State S2;
      // some uses
    });
    ```
 - Different memory device types possible as well

THE UNIVERSITY of EDINBURGH

epcc

DPC++

- Ensuring data types match across assignments helps performance

- Variable precision support

Data Type	Header File	Description
ac_int	<sycl/ext/intel/ac_types/ac_int.hpp>	Arbitrary-precision integer support
ac_fixed	<sycl/ext/intel/ac_types/ac_fixed.hpp>	Arbitrary-precision fixed-point number support
ac_fixed_math	<sycl/ext/intel/ac_types/ac_fixed_math.hpp>	Support for some non-standard math functions for arbitrary-precision fixed-point data types.
ac_complex	<sycl/ext/intel/ac_types/ac_complex.hpp>	Complex number support
ap_float	<sycl/ext/intel/ac_types/ap_float.hpp>	Arbitrary-precision floating-point number support
ap_float_math	<sycl/ext/intel/ac_types/ap_float_math.hpp>	Support for commonly used exponential, logarithmic, power, and trigonometric functions with ap_float type

THE UNIVERSITY of EDINBURGH

epcc

Performance issues

- Data dependencies can stall pipelines
  - Occupancy metric used to quantify how much of the hardware used
  - Aiming for 100% usage and 100% occupancy
- Full usage may lead to network issues
  - Reduced clock speed or stalls
- Automatic memory types may not be ideal
  - Chosen types can be tuned for performance
- Memory banking can also be tuned
- As can access coalescing
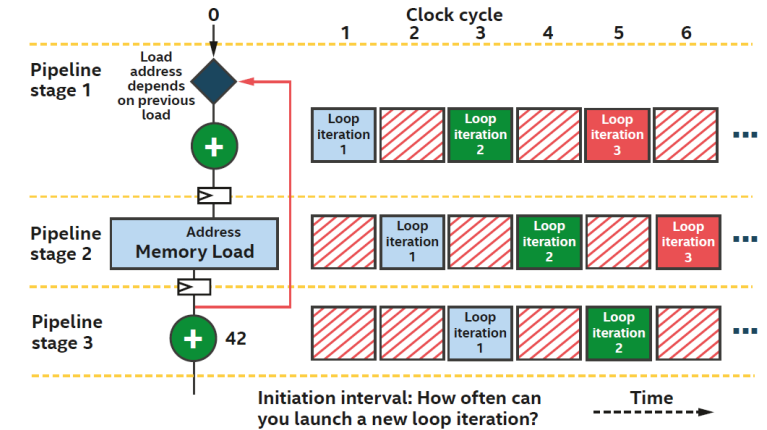
FPGA testbed

- EPCC is assembling an FPGA tested
  - Funded by the UKRI Excalibur programme
  - https://fpga.epcc.ed.ac.uk/

- Mixture of FPGAs to try
  - Intel Stratix...
  - Xilinx Alveo and Versal
  - Others to come

- Direct FPGA to FPGA networking

- Host resources and software environments

- Get in touch for access

THE UNIVERSITY of EDINBURGH

epcc